

SymEngine: A Fast Symbolic Manipulation Library

Ondřej Čertík, Isuru Fernando, Thilina Rathnayake, Abhinav
Agarwal, Sumith Kulal, Abinash Meher, Rajith Vidanaarachchi,
Shikhar Jaiswal, Ranjith Kumar

July 16, 2017

Outline

SymEngine

- ▶ Introduction
- ▶ Features
- ▶ Demo (Python, Ruby, Julia)
- ▶ Why C++, how to write safe code
- ▶ Internals of SymEngine
- ▶ Roadmap for using SymEngine in SymPy
- ▶ Benchmarks

Introduction

About SymEngine

- ▶ Symbolic manipulation library written in C++
- ▶ Thin wrappers to Python, Ruby, Julia, C and Haskell
- ▶ <https://github.com/symengine/symengine>
- ▶ <https://github.com/symengine/symengine.py>
- ▶ <https://github.com/symengine/symengine.rb>
- ▶ <https://github.com/symengine/symengine.jl>
- ▶ <https://github.com/symengine/symengine.hs>
- ▶ MIT licensed
- ▶ Started in 2012
- ▶ 46 contributors
- ▶ Runs on Linux (GCC, Clang, Intel), OS X (GCC, Clang), Windows (MSVC, MinGW, MinGW-w64)
- ▶ Part of the SymPy organization, but the C++ library is Python independent

Introduction

Goals

- ▶ Be the fastest symbolic manipulation library (open-source or commercial)
- ▶ Serve as the core for SymPy and Sage
- ▶ Serve as the default symbolic manipulation library in other languages thanks to thin wrappers (Python, Ruby, Julia, C and Haskell)

Choice of Language

Problem

SymPy speed is sometimes insufficient

- ▶ Handling of very large expressions
- ▶ Large calculations using small/medium size expressions

Let's fix that

- ▶ We tried: pure Python/PyPy, Cython, C, ...
- ▶ Investigated Julia, Rust, Scala, Javascript, ...
- ▶ Chose C++

Current Features

- ▶ Core (Symbols, +, -, *, /, **)
- ▶ Elementary Functions (sin, cos, gamma, erf)
- ▶ Differentiation, Substitution
- ▶ Matrices
- ▶ Polynomials (Piranha, Flint)
- ▶ Series Expansion
- ▶ Solvers
- ▶ Printing, Parsing
- ▶ Numeric Evaluation (Double/Arbitrary Precision)

Demo Time

Why Pure C++

- ▶ Fast in Release mode, but safe in Debug mode
- ▶ Compiler helps (not as good as Scala or Haskell, but much better than Python)
- ▶ Just one language to learn, thus easy to maintain (as opposed to several intertwined layers such as C + Cython + Python)
- ▶ Thin wrappers (that core developers do not need to maintain), all functionality in C++
- ▶ Easier to create bindings to other languages like Python, Julia, Ruby and Haskell

Why Pure C++: Fast in Release Mode

- ▶ Allows direct memory handling (allocation, deallocation, access)
- ▶ Allows to tweak how and when things are done
- ▶ It is possible to go to bare metal
- ▶ Allows reasonably high level abstractions (simple, maintainable code)

Why Pure C++: Safe in Debug Mode

- ▶ Reference counted pointers Teuchos::RCP (from Trilinos)
- ▶ Checks for dangling and null pointers (exception is raised)
- ▶ No raw pointers/references (use Ptr and RCP)
- ▶ Use a safe subset of C++
- ▶ Few other rules, e.g. how to use Ptr and RCP properly
- ▶ Possible to visually verify in a PR (pull request) review
- ▶ Hopefully eventually there are plugins to Clang to check automatically (since the rules are simple and static)
- ▶ As fast as raw pointers in Release mode (but it could segfault)

Conclusion: the code cannot segfault or have undefined behavior in Debug mode — always get an exception at runtime, or a compile error.

Internals of SymEngine

How Add, Mul and Pow work

- ▶ Add uses `std::unordered_map` (hashtable)
 - ▶ $2xy^2 + 3x^2y + 5 \rightarrow \{xy^2 : 2, x^2y : 3\}; coeff = 5$
- ▶ Mul uses `std::map` (red-black tree)
 - ▶ $2xy^2 \rightarrow \{x : 1, y : 2\}; coeff = 2$
- ▶ Pow just stores the base and exponent
- ▶ Each object is reference counted (RCP), very fast implementation in Release mode

Internals of SymEngine

Extensibility using Visitor Pattern

- ▶ All algorithms implemented using visitor pattern
- ▶ Algorithm is implemented in its own file, separate from the core
- ▶ Two virtual function calls (can be implemented in third party code or user code)
- ▶ Special version with just one virtual function call (faster, but must be compiled as part of the SymEngine source code)
- ▶ The speed difference between the two is minor

SymPy, SymEngine and the Interface

Using SymPy in SymEngine

- SymEngine will convert any SymPy object to a corresponding SymEngine object before doing any operation

```
>>> from symengine import symbols, Add
>>> import sympy
>>> x = symbols("x")
>>> y = sympy.symbols("y")
>>> x + y
x + y
>>> type(x+y)
<type 'symengine.lib.symengine_wrapper.Add'>
```

- What if there is no corresponding SymEngine object?

SymPy, SymEngine and the Interface

Using SymPy in SymEngine

- ▶ SymEngine will keep a reference to a SymPy object if there is no corresponding SymEngine object using Python/C API. SymEngine will use Python callbacks to evaluate the SymPy object

```
>>> e = x + sympy.Mod(x, 2)
>>> assert str(e) == "x + Mod(x, 2)"
>>> assert isinstance(e, Add)
```

```
>>> f = e.subs({x : 10})
>>> assert f == 10
```

```
>>> f = e.subs({x : 2})
>>> assert f == 2
```

SymPy, SymEngine and the Interface

Using SymEngine in SymPy

- ▶ `>>> from sympy.core.backend import symbols, sin, diff`
- ▶ Most things can be used unmodified
- ▶ Few things are fundamentally different (e.g. SymPy stores i as `ImaginaryUnit`, SymEngine has a `Complex` class)
- ▶ We will have a compatibility layer, probably similar to Python 2 and 3 support using the same source base.

Benchmarks

Benchmark setup

Benchmarks were run in a Intel(R) Core(TM) i5-5200U CPU @ 2.20GHz running Ubuntu 16.04 with gcc 5.4.0

- ▶ SymEngine master (with GMP and FLINT)
- ▶ GiNaC 1.6.6
- ▶ SymPy 1.0
- ▶ Mathematica 10.2.0.0
- ▶ Maple 2015.2

Benchmarks

Expand Benchmark


- ▶ $e = (x + y + z + w)^n$
- ▶ $f = e * (e + w)$
- ▶ Measure time taken for expanding f

- ▶ using SymEngine
using TimeIt

```
@vars x y z w
n = 30
e = (x + y + z + w)^n
f = e * (e + w)
@timeit expand(f)
```

Benchmarks

Expand Benchmark



Benchmarks

Modified GiNaC Benchmark

- ▶ Let e be the expanded sum of 2 symbols $\{a_0, a_1\}$ and $n - 2$ trigonometric functions $\{\sin(a_2), \sin(a_3) \dots \sin(a_{n-1})\}$ squared:
$$e \leftarrow (a_0 + a_1 + \sum_{i=2}^{n-1} \sin(a_i))^2$$
- ▶ Substitute $a_0 \leftarrow -\sum_{i=2}^{n-1} \sin(a_i)$
- ▶ Expand e again so it collapses to a_1^2

Benchmarks


Modified GiNaC Benchmark

```
from symengine import symbols, sin
from time import clock
n = 100
a0, a1 = symbols("a0, a1")
t = sum([sin(symbols("a%s" % i)) for i in range(2, n)])
e = a0 + a1 + t
f = -t

t1 = clock()
e = (e**2).expand()
e = e.xreplace({a0: f})
e = e.expand()
t2 = clock()
```

Benchmarks

Modified GiNaC Benchmark



Benchmarks


SymEngine Benchmark

- ▶ Series expansion of $\sin(\cos(x + 1))$ around $x = 0$
- ▶

```
RCP<const Symbol> x = symbol("x");  
int n = 15;  
RCP<const Basic> ex = sin(cos(add(integer(1), x)));  
auto t1 = std::chrono::high_resolution_clock::now();  
RCP<const Basic> res = series(ex, x, n);  
auto t2 = std::chrono::high_resolution_clock::now();
```


Benchmarks

SymEngine Benchmark



Benchmarks

PyDy Benchmark



Benchmarks


PyDy Benchmark

n	SymEngine + SymPy	SymPy only	Speedup
10	0.17 s	5.58 s	32.8x
15	0.44 s	22.07 s	50.1x
20	0.95 s	45.59 s	47.9x
30	3.11 s	162.80 s	52.3x
40	7.02 s	427.16 s	60.8x
50	13.95 s	915.83 s	65.6x
60	23.16 s	1596.37 s	68.9x

Table: Results

Benchmarks

PyDy Benchmark



Summary

- ▶ SymEngine aims to be the fastest C++ symbolic manipulation library
- ▶ Thin wrappers to other languages (Python, Ruby, Julia, C and Haskell)
- ▶ Easily usable as an optional backend in SymPy, Sage and PyDy

Thank You

GitHub:

- ▶ <https://github.com/symengine/symengine>
- ▶ <https://github.com/symengine/symengine.py>
- ▶ <https://github.com/symengine/symengine.rb>
- ▶ <https://github.com/symengine/symengine.jl>
- ▶ <https://github.com/symengine/symengine.hs>

Mailinglist:

- ▶ <http://groups.google.com/group/symengine>

Gitter:

- ▶ <https://gitter.im/symengine/symengine>